

Voice is context

How Cassini records your meetings into searchable memory

– Chris De Sousa



David Shrigley

1. Capture

- The recorder shouldn't hurt the call.

2. Transcription

- On your own host. No LLM in the pipeline.

3. Format

- A whole meeting in one file.

4. Intelligence

- No LLM, unless you say so.

Capture

The recorder shouldn't hurt the call.



(We spent most of those weeks on Google Meet)

The recorder shouldn't hurt the call





Cassini: save now, process afterward

DURING THE CALL

Talk call
everyone's
audio + video

Receive packets

Go + Pion (WebRTC),
plugged into the High
Performance Backend

Write to disk

one track per person,
into a multi-track MKV
- nothing decoded

AFTER

Decode & mix

decode the audio - mix the
tracks down to one

Transcribe

one pass per person -
names from signalling

Pack & publish

one Ogg Opus file → Nextcloud
Files (WebDAV)

Transcription

On your own host. No LLM in the pipeline.



names come from signalling,
not diarisation



Transcriber

Model · Parakeet - NVIDIA's
open speech-to-text model

Runtime · sherpa-onnx -
Kaldi's successor. Runs
inside our Go binary.



meeting.opus

an ordinary audio file

	MEMORY / PRICE	SPEED
Parakeet on CPU sherpa-onnx · 0.6B v3 · int8 — the default on a CPU	1.3 GiB RAM	30× real-time
Parakeet on CPU, full precision sherpa-onnx · 0.6B v3 · fp32 — same model as the GPU	2.9 GiB RAM	13× real-time
Parakeet on GPU sherpa-onnx · 0.6B v3 · fp32 · CUDA 12	3.4 GiB VRAM + 1.2 GiB RAM	258× real-time
MAI-Transcribe-2 Microsoft · hosted API, via OpenRouter	\$0.10 / audio hour intro price · audio leaves the building	149× real-time

× real-time = faster than the meeting itself: at 30× an hour of audio takes two minutes, at 258× about fourteen seconds · 30 minutes of our meetings, one machine (Ryzen 5 7600 + RTX 5060) · full report linked in the QR at the end


Format

A whole meeting in one file.

An audio file that
carries its own
transcripts

meeting.opus

Ogg · Opus codec



TITLE · DATE	The tags any player shows
TRANSCRIPT	Every word, with a time and a name
SPEAKERS	Who was in the room
PROVENANCE	Which engine, model and device made it
SUMMARY	If you asked for one
AUDIO_SHA256	A hash of the audio itself

Intelligence

No LLM, unless you say so.

Who is actually going to read all of these?

review... -29 · 56m	2026-03-23 · 38m	hiring.opus 2026-07-31 · 62m	2025-10-14 · 64m	Berlin talk.op... 2026-05-09 · 68m	2026-06-18 · 43m	microsite.opus 2026-03-24 · 26m	2026-07-27
.opus -18 · 30m	2026-06-09 · 21m	hiring.opus 2026-08-17 · 34m	Futto.opus 2026-05-31 · 68m	shaping – insi... 2025-10-18 · 60m	design review... 2025-10-24 · 21m	design review... 2026-06-03 · 62m	Marco – s 2025-12-29
– trans... -19 · 64m	2026-04-11 · 17m	planning.opus 2026-04-18 · 18m	Alex – the dec... 2025-12-27 · 38m	retro.opus 2026-03-20 · 19m	offsite prep.o... 2026-06-23 · 58m	hiring.opus 2025-11-26 · 18m	microsite 2026-01-21
call.op... -26 · 48m	2026-01-28 · 50m	standup.opus 2025-10-28 · 80m	support.opus 2025-11-11 · 68m	release 0.2.0... 2026-07-14 · 18m	offsite prep.o... 2025-10-02 · 39m	1:1.opus 2025-12-30 · 26m	retro.opu 2026-01-30
review... -29 · 56m	2026-03-23 · 38m	hiring.opus 2026-07-31 · 62m	client call.op... 2025-10-14 · 64m	Berlin talk.op... 2026-05-09 · 68m	ops sync.opus 2026-06-18 · 43m	microsite.opus 2026-03-24 · 26m	roadmap.o 2026-07-27

The default is to push a model in front of everything, whether you asked for one or not.

Ask Linear...

 Skills ▾

@

↑

The default - no model



Read it

the transcript,
with names



Search it

across all
your meetings



Tag it

meetings, and
moments inside them



The default - no model



Read it

the transcript,
with names



Search it

across all
your meetings



Tag it

meetings, and
moments inside them



With an LLM - opt-in



Summarise it

one meeting, or
a month of them



Act on it

decisions and
actions, with owners



Build on it

context for any
workflow or agent



Design review — 18 Aug

Design 2026-08-18 14:30 46m Ivan Chris

SUMMARY

The group worked through how the catalog should be organised once rooms exist, and agreed a flat list does not survive contact with more than a handful of meetings.

Decisions

- Selecting meetings into context is the affordance I'm least sure about.
- My worry is we end up with a chat box as the front door, which is exactly what we said we didn't want.

Follow-ups

- None.

INSIGHTS

Where the catalog design landed

2026-08-18 • Context from 4 meetings

Transcript

Select a word to start playing from there.

Ivan

0:00

I've got the branch open, so I'll share as we go.

Chris

0:00

Where the catalog design landed

Design

2026-08-18

anthropic/claude-sonnet-4.5

“What was decided, and what is still open?”

Context from 4 meetings


Design review — 18 Aug	2026-08-18
Design review — 18 Aug	2026-08-18
Design review — 17 Aug	2026-08-17
Design review — 17 Aug	2026-08-17

Across the design reviews the group converged on rooms as a grouping over the existing catalog, with the viewer reduced to a detail view. The flat list was accepted as the thing that breaks first.

What was decided


- Selecting meetings into context is the affordance I'm least sure about.
- My worry is we end up with a chat box as the front door, which is exactly what we said we didn't want.
- The catalog doesn't carry any notion of grouping, so everything lands in one flat list.

GoCassini > First test of Hermes/Cassini integration AUG 27

**Silvio Tomatis**

Summarize activity seen in meetings this week.

Ahem... @Hermes Summarize activity seen in meetings this

**Hermes**

This week (Aug 24–27) there are 6 recordings, all from the **Daily standup meeting** room. Let me pull the transcripts/summaries for — that downloads the full recording per meeting, so I'll run it in the background.

On it — this week (Aug 24–27) Cassini has **6 recordings**, all from the *Daily standup meeting* room:

Date	Time	Speakers	Duration
Mon Aug 24	11:13	5	~40 min
Tue Aug 25	11:13	5	~39 min
Wed Aug 26	10:38	1	~0.5 min (short stray recording)
Wed Aug 26	10:38	3	~35 min
Wed Aug 26	11:06	4	~27 min
Thu Aug 27	11:14	5	~39 min

I'm pulling the transcripts/summaries now (each requires downloading the full recording, so give me a couple of minutes). I'll post the summary here as soon as it's ready.

GoCassini > First test of Hermes/Cassini integration AUG 27

Meeting activity this week (Aug 24–27)

6 recordings, all **Daily Standup** — Mon–Thu, ~35–40 min each, plus a 35-second test clip. Attendees: Chris, Silvio, Ivan, Marco (chima), Alex.

Cassini — the dominant theme

Access control wrap-up (Mon–Thu)

- Ivan landed **PR 204** (room-ID backfill, merged) and drove **PR 208 / D594** — the fix for the ACL leak between opus upload and ACL application. ACLs now survive republish (frozen at first successful publish). Marked ready-for-review Thursday; merges if no objections.
- Wednesday had a dedicated pre-standup working session** (35 min, Ivan/Chris/Silvio) on access control. Key outcomes: principle that *attendees at the time of the meeting get access*; the dependency on `group_everyone` + team folder stands (Silvio challenged it, Ivan defended — shared ownership too hard otherwise); end-to-end leak tests deferred to maintenance mode; **opt-in publication is the only gap left before "Cassini stable"**.
- Ivan shaped the opt-in work in **D616** (agent is triaging it into spike/tasks); publish must keep working by default without the dependencies.

Insights / "Ask Your Meetings" (the next push)

- Chris reorganized Linear into an **"Ask Your Meetings + NCC Demo"** project and declared a **two week focus cycle** toward the Nextcloud

↑ FULL CONTROL · nothing leaves the building

Transcription on the infrastructure you own

CPU or GPU - whichever you have.
Still yours, even if someone else runs the box.

None

The transcript, search and tags are the whole product

A LLM model you run

An open-weight model on your own Ollama server · nothing leaves

A LLM model you pay for

Any OpenAI-compatible endpoint · your meetings leave the building

↓ FULL SERVICE · if you're fine trusting a third party

**Harness voice,
with full privacy.**



Cassini just now

The recording of **Voice is Context**
is ready to view.



voice-is-context-berlin.opus

audio · transcript · speakers · summary - one file



We're here all week, say hi!

gocassini.com/nextcloud-conf-2026